

## Lab 2: Profiler Report

The **Profiler system** is a tool designed to monitor and analyze the performance of various sections of code. It helps developers measure execution times, detect performance bottlenecks, and compare the performance of different code segments. The system supports both **hierarchical profiling** (nested code blocks) and **interleaved profiling** (sections that are exited before a previously entered section completes), making it a robust tool for detailed performance analysis.

### 1.1. Overall Architecture and Key Classes/Functions

The Profiler system is composed of several key components that work together to record and analyze the performance of code sections.

#### Key Components:

- **Profiler:** The core class responsible for managing profiling sections, calculating statistics, and generating reports.
- **ProfilerScopeObject:** A helper class that ensures sections are automatically started and exited using the RAII (Resource Acquisition Is Initialization) pattern.
- **TimeRecordStart and TimeRecordStop:** Classes used to record timestamps when sections are started and stopped. These use the user-defined class `time.hpp` for time measurement.
- **ProfilerStats:** Stores aggregated performance statistics for each profiled section, including total time, minimum time, maximum time, average time, and the number of calls.
- **Macros:** The `PROFILER_ENTER` and `PROFILER_EXIT` macros simplify the process of profiling sections of code for developers.

#### Key Classes and Functions:

##### 1. Profiler:

- `EnterSection(const char* sectionName):` Starts profiling a section and records the current time.
- `ExitSection(const char* sectionName, ...):` Stops the section profiling and calculates the elapsed time.
- `ReportSectionTime:` Records and updates the statistics with the elapsed time.
- `calculateStats():` Aggregates and computes performance statistics (e.g., total time, minimum/maximum time, average time).
- `printStats(), printStatsToCSV(), printStatsToJSON():` Outputs the profiling results to the console or files in CSV/JSON format.
- **Data Structures:**

- `map<char const*, ProfilerStats*> stats`: A map to store statistics for each section.
- `vector<TimeRecordStop> elapsedTimes`: Stores elapsed times for all profiled sections.
- `map<const char*, TimeRecordStart> sectionMap`: Tracks active sections.

## 2. **ProfilerScopeObject:**

- Automatically calls `EnterSection()` when created and `ExitSection()` when destroyed, ensuring that sections are correctly started and exited, even in case of exceptions or early function returns.

## 3. **TimeRecordStart and TimeRecordStop:**

- `TimeRecordStart`: Stores the start time of a section.
- `TimeRecordStop`: Stores the stop time and additional metadata (e.g., line number, file name, and function name).

## 4. **ProfilerStats:**

- Holds statistics for each profiled section, including the count (number of calls), total execution time, minimum time, maximum time, and average time. It also stores metadata such as the file name, function name, and line number where the section was executed.

## 1.2. Performance Considerations

The Profiler system is carefully designed to minimize the overhead of profiling itself, so it does not significantly affect the performance of the code being profiled.

### Key Performance Considerations:

- **Efficient Data Structures:** The system uses `std::map` and `std::vector` for efficient lookups, updates, and storage of performance data.
- **Pre-allocation of Vectors:** To minimize memory allocation overhead during profiling, vectors for storing timestamps (`TimeRecordStart` and `TimeRecordStop`) are pre-allocated based on estimated usage.
- **Minimized System Calls:** The `GetCurrentTimeSeconds()` function uses lightweight time measurement techniques (such as `std::chrono`) to reduce the performance impact of time collection.
- **Optimized Sorting Algorithms:** Examples like early exit conditions in sorting algorithms (e.g., Insertion Sort) illustrate how profiling can guide developers to improve algorithmic performance.

By focusing on these considerations, the Profiler ensures that it can be applied both to small, high-frequency code sections and long-running algorithms without introducing significant performance overhead.

### 1.3. Example-Driven Guide to Using the Profiler API

The Profiler API is designed to be straightforward and easy to integrate into any codebase. Developers can use the **PROFILER\_ENTER** and **PROFILER\_EXIT** macros to profile any section of code.

#### Example 1: Simple Profiling of a Code Section

```
#include "profiler.hpp"

void someFunction() {
    PROFILER_ENTER("Some Function");
    // Perform some task
    for (int i = 0; i < 1000; i++) {
        // Code to be Profiled
    }
    PROFILER_EXIT("Some Function");
}
```

In this example, the **PROFILER\_ENTER** macro starts profiling the section of code named "Some Function". Once the execution reaches the **PROFILER\_EXIT** macro, the profiler will stop tracking and record the elapsed time for the section. This allows developers to see the execution time for "Some Function" and include it in the profiler's output report.

#### Important: Managing the Profiler Singleton

When using the Profiler, it's important to manage the **Profiler** instance properly to ensure efficient memory usage and avoid resource leaks. The **Profiler** is implemented as a **singleton** to ensure that there is only one instance running at any given time across the application.

1. **Getting the Profiler instance:** The profiler should be initialized once in the program and set to nullptr at the end.
2. **Resource Cleanup:** Once profiling is complete, it is essential to free up resources by deleting the Profiler instance and setting it to nullptr.

### 1.4. Hierarchical and Interleaved Profiling

The Profiler system is designed to handle both **hierarchical** and **interleaved profiling**. These advanced profiling techniques allow developers to track complex code structures where sections

of code are either nested within each other (hierarchical profiling) or overlap in execution (interleaved profiling). This functionality is critical for capturing the performance of complex, real-world applications.

### **Hierarchical Profiling (Nested Sections):**

Hierarchical profiling enables developers to track code sections that are nested within one another. This is common when one function calls another, and both functions need to be profiled independently. The Profiler handles this by maintaining a map of active sections, ensuring that each section's start and end times are tracked without confusion.

- **How It Works:** Each time a section is entered, the Profiler stores its start time in the `sectionMap`. When a section is exited, its start time is retrieved from the map, and the elapsed time is calculated. This design allows the Profiler to accurately manage nested sections, ensuring that the start and stop times of each section are preserved, even when sections are deeply nested.

### **Interleaved Profiling (Overlapping Sections):**

Interleaved profiling is essential when sections of code overlap in execution and are not exited in the same order they were entered. This occurs in scenarios where a developer may need to exit a section before another previously entered section completes, creating a "crossed" execution flow.

- **How It Works:** The Profiler handles interleaved profiling using a map (`sectionMap`) to store the start time of each section. When a section is exited, the system retrieves its start time from the map, calculates the elapsed time, and removes the entry, allowing sections to be stopped in any order.

### **Conclusion:**

The **Profiler system** is a flexible and efficient tool for monitoring the performance of code sections. Its support for hierarchical and interleaved profiling makes it highly versatile for complex applications. By providing a simple API (`PROFILER_ENTER` and `PROFILER_EXIT`), developers can easily integrate it into their projects to gain detailed insights into the performance of various sections of their code.

Performance considerations such as efficient data structures, pre-allocation, and minimized system calls ensure that the Profiler imposes minimal overhead. By tracking execution times and bottlenecks, developers can use the Profiler to optimize their code, improve efficiency, and reduce execution time.

## **Section 2: Visualization**

The Profiler system is designed to efficiently capture and analyze performance statistics for various sections of code. To enhance its functionality, a visualization tool was implemented, allowing users to view the profiling data in an interactive and graphical format. This tool

leverages the **Chart.js** library to create multiple visualizations from the collected data, which is stored in **CSV** and **JSON** formats.

The **Profiler** class in the system is responsible for tracking performance data, calculating statistics, and outputting these results to both CSV and JSON files. The decision to choose these two formats was based on the following considerations:

- **CSV (Comma-Separated Values):** CSV is a lightweight, plain-text format that is widely used for data processing and analysis. It is human-readable and can be easily opened in spreadsheets or imported into data visualization tools for further processing. Its structure makes it straightforward to parse and analyze, especially when dealing with numerical data like execution times.
- **JSON (JavaScript Object Notation):** JSON provides a more structured format for representing hierarchical data. In addition to being easily readable by humans, JSON is widely used in web applications and APIs for transmitting data. The hierarchical structure of JSON makes it well-suited for representing complex datasets, such as profiling results that include nested sections of code. JSON is particularly advantageous when building interactive web-based visualizations, as it can be seamlessly integrated with JavaScript frameworks.

By outputting the profiling data in both **CSV** and **JSON**, the system provides flexibility, allowing users to choose the format that best suits their needs. CSV is ideal for quick data analysis and simple tools, while JSON is preferred for web-based applications and structured data representations.

## 2.1 Justification of Visualization Chart.js

The choice of **Chart.js** as the visualization technology for this profiler was driven by several factors that align with the goals of ease of use, interactivity, flexibility, and responsiveness. Below are the key reasons why **Chart.js** was chosen:

1. **Simplicity and Ease of Integration:**
  - **Chart.js** is lightweight and easy to integrate with web-based applications. It provides simple APIs that allow for quick setup and creation of visualizations. Given the requirement to output data in a user-friendly, interactive format, Chart.js offers a straightforward solution that works seamlessly with JavaScript, making it a perfect choice for visualizing data stored in **JSON**.
2. **Cross-Browser Compatibility:**
  - One of the key advantages of **Chart.js** is its ability to work across all modern browsers, ensuring that the visualizations can be viewed consistently by users regardless of their platform. This ensures the profiler's visualizations can be accessed by developers and users on any system, from Windows to macOS and Linux, without any issues.

### 3. Interactive and Dynamic Visualizations:

- **Chart.js** supports dynamic and interactive charts, allowing users to hover over data points to see detailed tooltips or zoom into specific sections of the data. This aligns with the need to provide insights into the profiled code sections, offering interactive feedback as developers explore different sorting algorithms and compare performance across runs.
- The tool also supports a variety of chart types (bar, line, pie, etc.), which allows for flexibility in representing different kinds of performance data, such as runtimes, average times, and comparisons between different algorithms.


### 4. Customizability and Responsiveness:

- **Chart.js** is highly customizable, enabling users to modify the appearance, color schemes, tooltips, and layout to meet specific needs. In this profiler system, custom colors and tooltips were used to make the charts easier to interpret.
- It also provides responsive layouts, which means the visualizations adjust dynamically to different screen sizes, making them suitable for both desktop and mobile viewing. This makes the profiler's output accessible across a variety of devices.

## 2.2 Visualization Dashboard

**Figure 1: Sorting Algorithm Performance Comparison**

Sorting Algorithm Performance Comparison



#### Performance Analysis:

**Fastest:** Optimized Insertion Sort4 - Early Exit (0.000025 seconds)

**Slowest:** Optimized Insertion Sort3 - Binary Search (0.056718 seconds)

Speed Difference: 2268.71x

Number of algorithms compared: 4

This chart, implemented as a horizontal bar chart, displaying the **total runtime** of different sorting algorithms to measure the differences between the 4 implementations of a sorting algorithm. This functionality can be applied to other programs to measure the total runtime of different algorithms and coding blocks.

The **x-axis** represents the total runtime in seconds, while the **y-axis** lists the different sorting algorithms being compared. Each bar represents the total time it took to execute the respective sorting algorithm, which is directly drawn from the profiling data. Users are able to hover over each bar chart and find the total runtime of the algorithms.

**Importance:** This chart allows for a quick comparison of the overall performance of each sorting algorithm. By visualizing the total runtime, it becomes clear which optimization techniques yield the most significant improvements in performance. The color-coded bars help identify the fastest and slowest algorithms, allowing developers to see which optimization strategies are most effective.


Below the sorting algorithm performance chart is a **Performance Analysis Panel** that provides a textual summary of key performance statistics. This includes:

- The **fastest algorithm**, with its total runtime.
- The **slowest algorithm**, with its total runtime.
- The **speed difference** between the fastest and slowest algorithms.
- The **number of algorithms** being compared.

**Importance:** This panel gives a concise summary of the performance data, allowing developers to quickly interpret which algorithm is most efficient and how much faster or slower the alternatives are. The use of color (green for fastest, red for slowest) visually reinforces these key points.

**Figure 2 Section Analysis by Function**






In this section, a user can select a specific **function** from a dropdown menu to analyze the performance of different subsections within that function. The chart that follows shows multiple metrics for each section, including total time, average time, min time, and max time.

The **y-axis** lists the sections of the selected function, while the **x-axis** displays time values for each of the metrics. Users are able to see the exact time in seconds by hovering over each respective bar chart.

**Importance:** This visualization is crucial for analyzing individual sections of the code and identifying bottlenecks within specific functions. By comparing metrics such as average time and maximum time, developers can pinpoint areas of the code that might require further optimization, especially for functions that are repeatedly called or executed in nested loops.

**Figure 3 - Performance Improvement Trend**

Performance Improvement Trend



Optimized Insertion Sort2 - Shifting: ↓ 1.92%

Optimized Insertion Sort3 - Binary Search: ↑ 26.42%

Optimized Insertion Sort4 - Early Exit: ↓ 99.96%



This line chart illustrates the **performance improvement trend** across different optimized sorting algorithms, compared to the baseline. It shows how the total runtime decreases as each optimization is applied:

- Baseline Insertion Sort
- Optimized Insertion Sort2 - Shifting
- Optimized Insertion Sort3 - Binary Search
- Optimized Insertion Sort4 - Early Exit

The **x-axis** lists the sorting algorithms, and the **y-axis** displays the total runtime in seconds.

Additionally, a **percentage change panel** is included below the chart, showing the relative performance improvement (or degradation) between each successive algorithm.

**Importance:** This trend chart provides a clear visual representation of how each optimization affects performance. Developers can easily track the cumulative effect of various optimizations, helping them understand the performance gains achieved with each technique. The percentage change panel offers further insight into the relative improvement, allowing for more precise performance analysis.

### 3.0 Comprehensive Analysis

#### 3.1. Initial Algorithm or Operation Being Profiled:

The algorithm being profiled is **Insertion Sort**, which is a simple sorting algorithm that builds the final sorted array one item at a time. It operates by iterating through the list, comparing each element with its predecessor, and shifting the predecessor elements to the right until it finds the correct position for the current element. It is most efficient for small datasets or nearly sorted data but has a time complexity of  $O(n^2)$  in the average and worst cases, making it inefficient for large datasets.

##### Basic Steps of Insertion Sort:

- Iterate through the array from the second element to the end.
- For each element, compare it with the elements before it.
- Shift all elements larger than the current element to the right.
- Insert the current element at the correct position.

In this profiled implementation, the algorithm performs these steps on a large dataset of random integers to assess its performance.

#### 3.2. Performance Modifications:

##### Modification 1: Optimized Shifting

- **Rationale:** The first modification aims to optimize the shifting process in the insertion sort by breaking the key selection and shifting into smaller steps. The goal is to separate these two logical steps and better understand the performance of each part of the algorithm.
- **Details:**
  - The algorithm still follows the basic insertion sort principles.
  - However, by profiling the "key selection" and "element shifting" separately, the developer can isolate and analyze the time spent in each part.
  - While this does not change the overall complexity, it provides more granular insights into which part of the algorithm is taking the most time.

### Modification 2: Binary Search for Key Placement

- **Rationale:** The second modification introduces **binary search** for finding the correct location to insert the current element. Instead of scanning all previous elements linearly, binary search reduces the number of comparisons to  $O(\log n)$  by exploiting the fact that the array is partially sorted during each pass.
- **Details:**
  - Binary search is used to locate the correct position for the key element.
  - Once the correct position is found, the elements between that position and the current index are shifted as in the original insertion sort.
  - This change reduces the number of comparisons during each insertion but keeps the shifting process, so the time complexity remains  $O(n^2)$  due to the shifts.

### Modification 3: Early Exit Optimization

- **Rationale:** This modification introduces an **early exit** condition, where the algorithm terminates if the array becomes sorted before reaching the last element. This optimization is beneficial in cases where the array is already sorted or nearly sorted, as it avoids unnecessary comparisons and shifts.
- **Details:**
  - During each pass, the algorithm tracks whether any shifting occurred.
  - If the algorithm detects that no shifting is needed (i.e., the array is already sorted), it exits early without completing the remaining iterations.
  - This significantly improves performance in best-case scenarios, where the array is either fully sorted or nearly sorted.

These modifications focus on improving the overall efficiency of the insertion sort algorithm and isolating key sections of the code to provide more granular performance data. The optimizations

are not expected to change the worst-case time complexity of  $O(n^2)$  but can improve the performance for specific datasets or under certain conditions (e.g., nearly sorted arrays). By implementing and profiling these changes, the impact of each modification on runtime and computational cost can be analyzed in detail.

### 3.3 Performance Analysis Demonstration

#### Baseline Insertion Sort

- **Section Name:** Baseline Insertion Sort
- **Total Time:** 0.0457443 seconds
- **Key Sections:**
  - Outer Loop: 0.0457358 seconds
  - Key Selection: 0.0438207 seconds
  - Element Shifting: 0.0402486 seconds
- **Analysis:** In the baseline implementation, the algorithm performs the traditional insertion sort where each key is selected, and elements are shifted iteratively. Most of the time is spent in Element Shifting, which is expected since the algorithm shifts all elements greater than the key until the correct position is found. The Key Selection step is relatively fast as it only selects the next unsorted element.

#### Optimized Insertion Sort 2 - Shifting

- **Section Name:** Optimized Insertion Sort2 - Shifting
- **Total Time:** 0.0448659 seconds (1.92% improvement from baseline)
- **Key Sections:**
  - Outer Loop: 0.0448624 seconds
  - Key Selection: 0.043098 seconds
  - Element Shifting: 0.0367443 seconds (8.94% improvement from baseline)
  - Key Placement: 0.0013562 seconds
- **Analysis:** In this optimization, the insertion sort is slightly modified to profile the **Key Placement** as a separate operation, which helped reduce the time spent in the Element Shifting section. By focusing on minimizing unnecessary shifts, we observed a 1.92% reduction in the total time. While the reduction is modest, it highlights that the Key Placement operation is fast and that reducing unnecessary shifts improves performance incrementally.

#### Optimized Insertion Sort 3 - Binary Search

- **Section Name:** Optimized Insertion Sort3 - Binary Search
- **Total Time:** 0.0567178 seconds (26.42% increase from previous)
- **Key Sections:**
  - Outer Loop: 0.0567137 seconds
  - Key Selection: 0.0547561 seconds
  - Binary Search: 0.0019664 seconds
  - Element Shifting: 0.0472606 seconds (increase from previous shifting time)
- **Analysis:** Although the introduction of **Binary Search** was expected to optimize key placement, it surprisingly increased the total time by 26.42%. The reason lies in the overhead introduced by binary search and its inability to significantly reduce the number of shifts required (shifting is still  $O(n)$  for each element). The performance improvement in Key Selection was overshadowed by the time spent in Element Shifting, which increased compared to the previous version. This suggests that while binary search minimizes comparisons, the cost of shifting remains the dominant factor in the algorithm's overall runtime.

#### Optimized Insertion Sort 4 - Early Exit

- **Section Name:** Optimized Insertion Sort4 - Early Exit
- **Total Time:** 0.000025 seconds (99.96% improvement)
- **Key Sections:**
  - Outer Loop: 0.0000239 seconds
  - Key Selection: 0.0000016 seconds
  - Element Shifting: 0.0000002 seconds
  - Early Exit Check: 0.0000137 seconds
- **Analysis:** This optimization introduces an **early exit** condition, which drastically improves performance. The total time decreased by a staggering 99.96%, which shows the effectiveness of this technique for nearly sorted or fully sorted datasets. Since the array was already sorted early on in the process, the algorithm exited before completing all iterations, resulting in minimal time spent on shifting or selecting elements. This result demonstrates the importance of adapting sorting algorithms based on the characteristics of the input data.

The **Baseline Insertion Sort** serves as the starting point for analyzing performance. Modifications made to improve performance include separating the shifting and key placement steps, using binary search for key placement, and introducing an early exit condition.

- The **Shifting Optimization** yielded a small improvement (1.92%), demonstrating that optimizing element shifting reduces the time spent on unnecessary operations.
- **Binary Search** introduced overhead, leading to a 26.42% increase in total time. Although binary search minimizes comparisons, the time spent on shifting remains significant.
- The **Early Exit Optimization** proved to be the most impactful, reducing the total time by 99.96%. This optimization highlights the value of exiting early when the array is already sorted, eliminating the need for unnecessary iterations. This algorithm was 2268.71x faster than the slowest insertion sort algorithm that utilized binary search.

By using the profiler's **Section Analysis by Function**, we were able to isolate key sections of the algorithm and measure how each modification influenced performance. This detailed analysis allows developers to make informed decisions about when and how to apply optimizations depending on the nature of the data being sorted.